



(a) Fleet: one teacher, one dataset, one loss; only capacity varies



(b) Dispatch: the smallest student is its own router



(c) Two measured operating points on the identical 872 SST-2 prompts

Efficiency: $s_1 \rightarrow s_4$, $t = 0.98$

-87.3% energy (1,275.2 mJ)
89.68% accuracy (-3.78 pts vs T)

Matching: $s_2 \rightarrow T$, t calibrated on $m = 654$ held-out

-23.8% energy
-0.07 $\pm$ 0.09 pts, indistinguishable from T

Oracle over the fleet (perfect gate): 869.3 mJ = 8.7% of T at 97.48% $\Rightarrow$ 91.3% of T 's energy is recoverable

Ensemble of all four: 5,414.6 mJ = 54% of T before any escalation, errors 3.2-7.6 $\times$ correlated $\Rightarrow$ closed